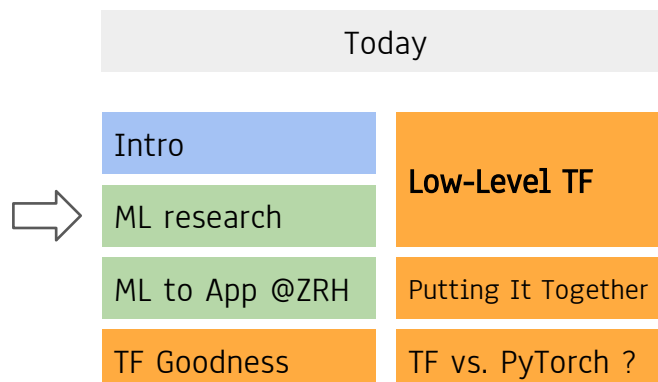




# From PyTorch to TensorFlow

EE-559 Deep Learning, EPFL, 5/23/18  
Andreas Steiner, Guest Lecture  
<https://fleuret.org/dlc>

## Agenda



# Diabetic Retinopathy



(Gulshan 2016)



# Pathology: Tumor Detection

- Inception v3 network, sliding window
- 242 slides -  $10^7$  299x299 patches
- 82.7% sensitivity (@7 FP/slide) vs. 73.2% pathologist



(Liu 2017)

# Pathology : Augmented Reality



(Po-Hsuan 2018)


# Transcribing Medical Records

- ~6/11h workday spent on electronic health records
- 14,000 hours anonymized conversations
- Attention model : 18% word error rate

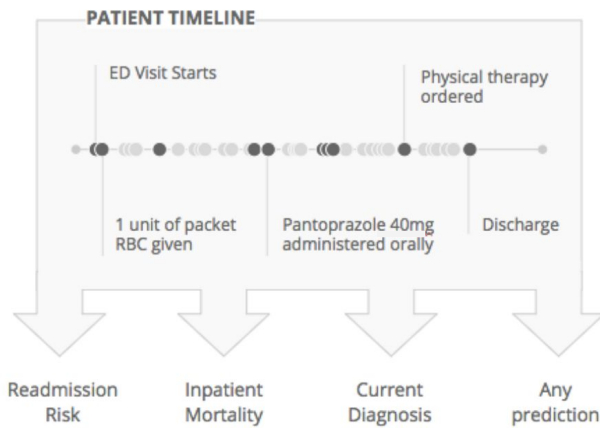


(Chiu 2017)  
(Chan 2015)

$\langle \text{sos} \rangle y_1, y_2, \dots, y_5 \langle \text{eos} \rangle$



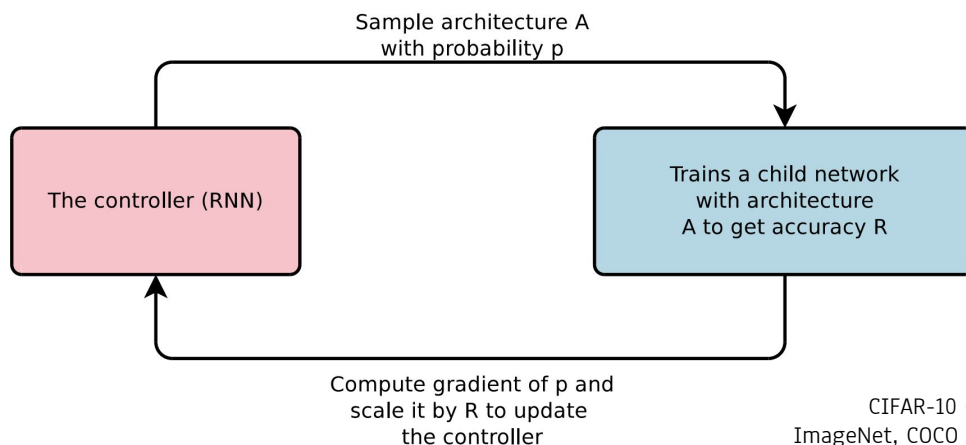
# Medical Records



- Predictive modeling from EHR
- Deep learning from raw FHIR data
- 216k patients  
46,864M datapoints
- Outperforms state-of-the-art in all predictions

(Rajkumar 2018)

# AutoML



CIFAR-10 (Zoph 2016)  
ImageNet, COCO (Zoph 2017)  
Optimizers (Bello 2017)  
Activation Functions (Ramachandran 2017)

# Datasets

## Youtube 8M

<https://research.google.com/youtube8m/index.html>

7M videos, 4716 classes

## Youtube BB

<https://research.google.com/youtube-bb/>

240k videos, 11M human annotations

## Atomic Visual Actions

<https://research.google.com/ava/>

58k segments, 210k action labels, 80 actions

## AudioSet

<https://research.google.com/audioset/>

2.1M human-labeled 10s sound clips, 632 classes

## Speech Commands

<https://ai.googleblog.com/2017/08/launching-speech-commands-dataset.html>

65k utterances, 30 words

## OpenImages V4

<https://storage.googleapis.com/openimages/web/>

9.2M images, 30M human-verified labels, 20k classes

1.9M images, 15M bounding boxes, 600 classes





Figure derived from Chan, et al. 2015

<https://distill.pub/2016/augmented-rnns/>  
(Olah 2016)



MIXED3A

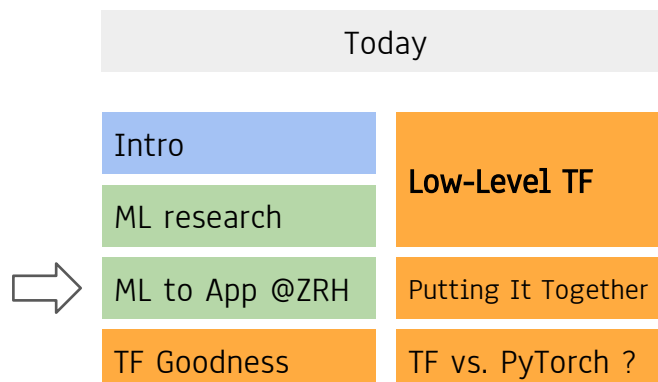
<https://distill.pub/2018/building-blocks/>  
(Olah 2018)

# AI Residency

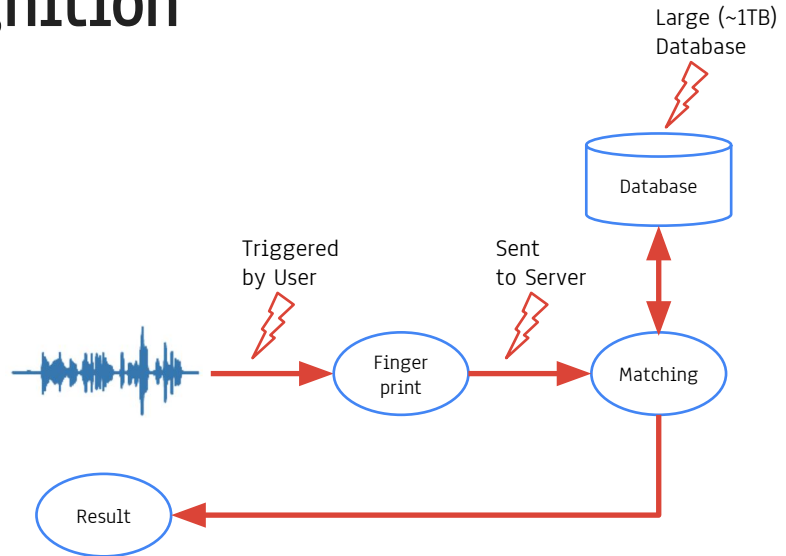
Ad [g.co/airesidency](https://g.co/airesidency)

- 12-month research training role
- For whom?  
Candidates with BSc/MSc/PhD in STEM field  
... or any background with passion for ML.
- Applications for 2019 will open in Fall 2018  
[g.co/airesidencysignup](https://g.co/airesidencysignup)

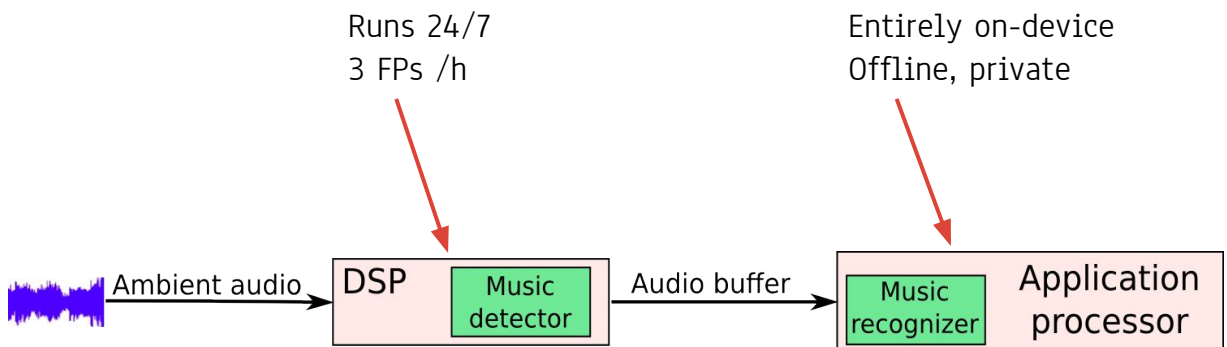
## Agenda



# Music Recognition




# Music Recognition 2.0



(Agüera y Arcas 2017)

# Music Detector



- 6 convolutional layers (stride=2) - prediction every 0.64s
- 8k parameters - 10k memory footprint (quantization!)
- Trained on AudioSet data

(Agüera y Arcas 2017)

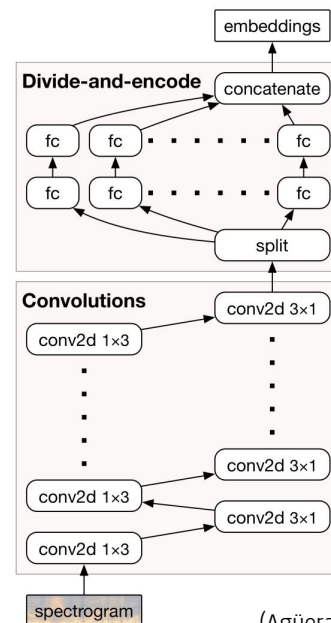
# Music Recognizer

Neural Network Fingerprinter

- Same input as Music Detector
- Triplet loss, 96-dim embeddings

Reduced size by:

- Separable convolutions
- FC : divide-and-conquer
- Quantization



(Agüera y Arcas 2017)



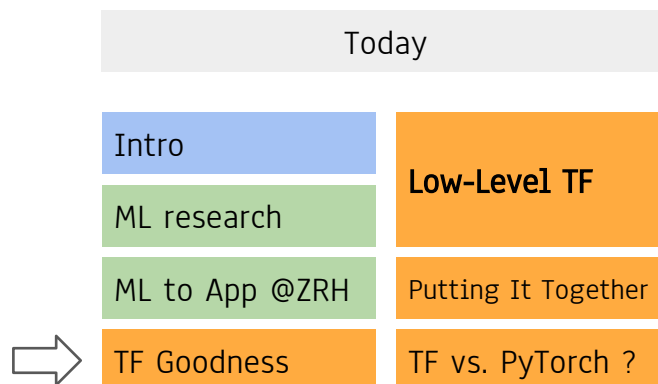
# Now Playing (on Pixel 2)

10k songs, <1% daily battery.


- On-device ML, private, always on.
- Requires tricks for reducing memory / power (cascading models, optimized models, hardware)



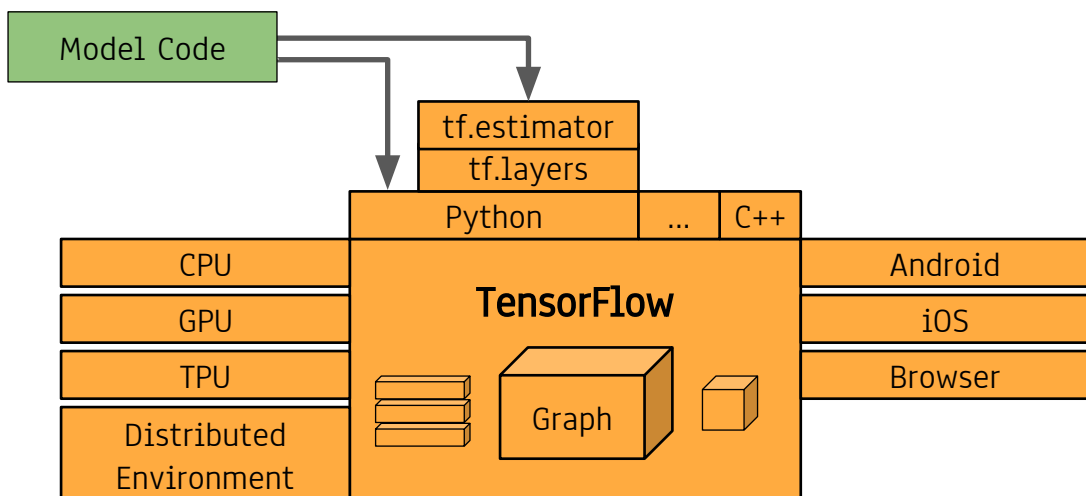
## Agenda



# Stars on Github

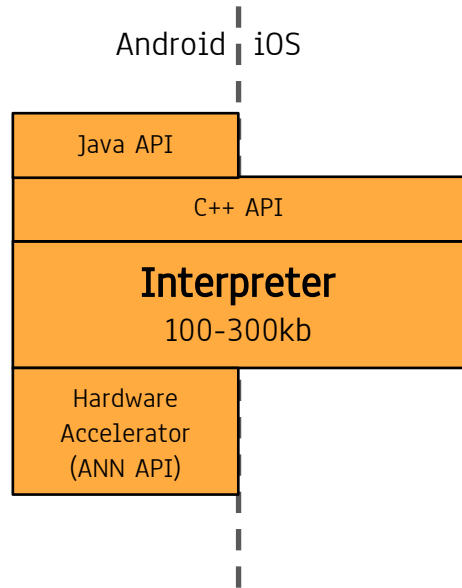


# Dealing With Complexity




# TensorFlow Lite

- Lightweight
- Low latency
- Operators (float/quantized) tuned for mobile performance
- Convert existing models to .tflite format



# TensorFlow.js



- Supports 90+ Ops, 32+ layers
- Train / run in browser
- Import SavedModel/Keras models
- ~1.5-2x slower than Python/AVX

Try it out:

<https://js.tensorflow.org>

# TensorFlow-Hub

```
import tensorflow_hub as hub

embed = hub.Module("https://tfhub.dev/google/"
                   "universal-sentence-encoder/1")

embedding = embed([
    "The quick brown fox jumps over the lazy dog."])
```

(Cer 2018)

- Module : self-contained piece of TF graph including weights and assets.
- Canonical URLs for versioned modules.
- For direct application, transfer learning, ...
- **Beware** of model bias !

Many more models under [github.com/tensorflow/models](https://github.com/tensorflow/models)


# Tensor Processing Units (TPUs)



(Jouppi 2017)

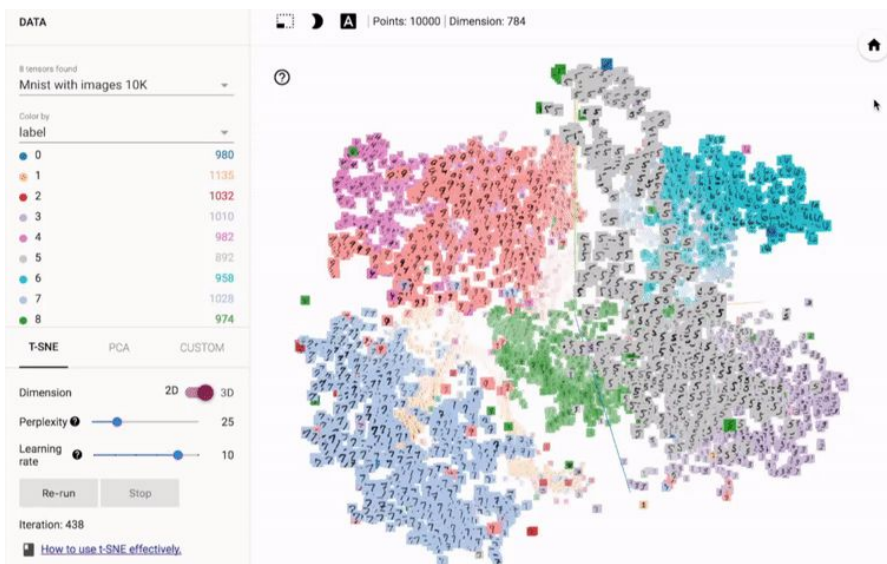
180 TFLOPS, 64 GB HBM memory, 2400 GB/s BW - [g.co/tpusignup](https://g.co/tpusignup)

# Distributed TensorFlow

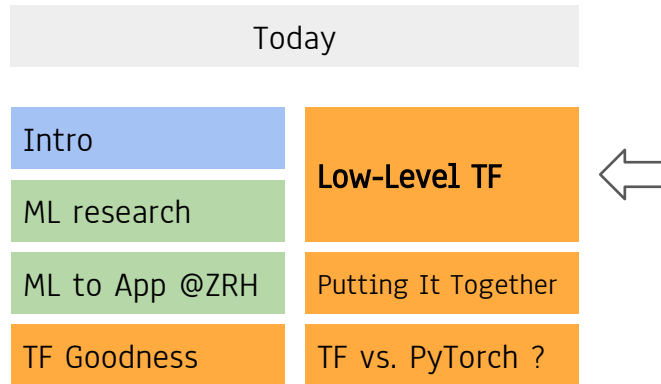


- Lot of flexibility.
- Communication within process / over network.
- Graph can be split or replicated.
- Easy to use with higher level API.

# TensorBoard



# Agenda



## Low-Level TensorFlow – Overview

- The Graph
- The Session
- Shapes
- Variables

Code

```
import tensorflow as tf  
print(tf.__version__)
```

1.7.0

Output

Find all code from slides here:


<https://goo.gl/EKVKWz>

# TensorFlow Versions

Semantic Versioning : MAJOR.MINOR.PATCH

<https://github.com/tensorflow/tensorflow/releases>


<https://github.com/tensorflow/tensorflow/blob/master/RELEASE.md>



## The Graph\*

**Ops** - Nodes in the graph. Perform computation on Tensors

**Tensors** - Edges in the graph. Input/output of Ops




\* always present, but might be invisible

# The Graph

**Ops** - Nodes in the graph. Perform computation on Tensors


**Tensors** - Edges in the graph. Input/output of Ops



```
x = tf.constant(1.0, name="x")
add = tf.add(x, 2.0, name="add")
mul = tf.multiply(add, 3.0,
                  name="mul")
```

## Naming, Placement, Scope

```
c = tf.constant(0, name="c")
c_1 = tf.constant(2, name="c_1")
with tf.device("/device:GPU:0"):
    with tf.name_scope("sum"):
        s1 = tf.constant(0, name="s1")
        with tf.name_scope("prod"):
            p1 = tf.constant(1, name="p1")
            p2 = tf.constant(2, name="p2")
            p12 = tf.multiply(p1, p2,
                             name="mult")
            sp = tf.add(s1, p12, name="add")
stack = tf.stack([sp, c],
                 name="stack")
```





# Conversion, Overloading, Broadcasting

```
a = tf.constant([1, 1])
c = a + 2
```

Conversion  
2 → tf.constant(2)

Overloading  
a + 2 → tf.add(a, 2)

Broadcasting  
2 → [2, 2]

```
a = tf.constant([1, 1])
c = tf.add(a, tf.constant([2, 2]))
```

## The Graph – Inspection

```
graph.get_operations()
```

```
[<tf.Operation 'Const' type=Const>,
 <tf.Operation 'add/y' type=Const>,
 <tf.Operation 'add' type=Add>]
```

```
graph.get_operation_by_name('Const')
```

```
<tf.Operation 'Const' type=Const>
```

```
graph.get_tensor_by_name('Const:0')
```

```
<tf.Tensor 'Const:0' shape=(2,) dtype=int32>
```

(helper function  
practical session)

```
show_graph(graph)
```

```
add/y
Operation: Const


Attributes (2)
dtype      {"type":"DT_INT32"}
value      {"tensor":
            {"dtype":"DT_INT32","tensor_s
            hape":{},"int_val":2}}

Inputs (0)
Outputs (0)
```



# The Session – Makes the Tensors Flow

The Session **executes the Graph**, similar to the JVM running Java bytecode.



build graph...

```
x = tf.constant(1.0, name="x")
add = tf.add(x, 2.0, name="add")
mul = tf.multiply(add, 3.0, name="mul")
```

...run graph

```
with tf.Session() as sess:
    mul_, = sess.run([mul])
    print(mul_)
```

## (Default) Graph, (Interactive) Session

```
tf.reset_default_graph()
x = tf.constant(3.1415)
sess = tf.InteractiveSession()
print(x.eval())
```

Start with a clean slate.

`x.eval()` is shortcut for `sess.run([x])`

```
graph = tf.Graph()
with graph.as_default():
    x = tf.constant(2.718)
with tf.Session(graph=graph) as sess:
    print(x.eval())
```

New Ops always attached to default graph.

Provide graph as argument.

```
with tf.Graph().as_default(), \
    tf.Session() as sess:
```

```
x = tf.constant(1.414)
print(sess.run([x]))
```

Slides often only show this code.

# Placeholders, Feeding

```
x = tf.placeholder(tf.float32, shape=[2], name="x")
norm = tf.norm(x)
```

```
norm.eval(feed_dict={x: [3, 4]})
```

```
5.0
```

```
sess.run([x, norm], feed_dict={"x:0": [3, 4]})
```


```
[array([3., 4.], dtype=float32), 5.0]
```

```
sess.run([x, norm])
```

```
InvalidArgumentError ... must feed a value ... tensor 'x'
```

```
sess.run([x, norm], feed_dict={x: [3, 4], norm: 0})
```

```
[array([3., 4.], dtype=float32), array(0., dtype=float32)]
```



# Graph/Session Quiz

```
with tf.Graph().as_default() as graph:
    x = tf.constant(1)
    x2 = tf.multiply(x, x)
```

```
with tf.Session() as sess:
    print(sess.run(x2))
```

```
ValueError: Tensor ... is not an element of this graph.
```

```
with tf.Graph().as_default() as graph:
    x = tf.constant(1)
    x2 = tf.multiply(x, x)
```

```
with tf.Session(graph=graph) as sess:
    print(sess.run(x2))
```

```
with tf.Graph().as_default():
    x = tf.constant(1)
    x2 = tf.multiply(x, x)
```

```
with tf.Session() as sess:
    print(sess.run(x2))
```

# Providing Functions, Not Tensors

```
inputs = get_inputs_tensor()
labels = get_labels_tensor()

model.train(inputs, labels)
```

⇒ These tensors are already attached to a graph.

```
def input_fn():
    inputs = get_inputs_tensor()
    labels = get_labels_tensor()
    return inputs_labels

model.train(input_fn)
```

⇒ `model.train()` can decide which graph tensors are attached to.

## The Shapes\*

1      **scalar**  
rank=0,  
shape=()

[1, 2]      **vector**  
rank=1  
shape=(2,)

[[1, 2, 3],  
 [2, 4, 6]]      **matrix**  
rank=2  
shape=(2, 3)

[[[1, 2],  
 [3, 4]],  
 [[5, 6],  
 [7, 8]]]      **tensor**  
rank=3  
shape=(2, 2, 2)

```
x = tf.random_normal(shape=(2, 2))
x.shape
```

```
TensorShape([Dimension(2), Dimension(2)])
```

```
x.shape.as_list()
```

```
[2, 2]
```

```
tf.reshape(x, (1, -1)).shape
```

```
TensorShape([Dimension(1), Dimension(4)])
```

\* the cause of, and solution to, all TensorFlow crashes!

# Partially Known Shapes

```
x = tf.placeholder(tf.float32, [None, 4])  
x.shape
```

```
TensorShape([Dimension(None), Dimension(4)])
```

```
x2 = tf.reshape(x, (2, -1))  
x2.eval({x: [range(4)]})
```

```
array([[0., 1.], [2., 3.]], dtype=float32)
```

```
batch_size = tf.shape(x)[0]  
x2b = tf.reshape(x, [batch_size, 2, -1])  
x2b.eval({x: [range(4), range(0, 8, 2)]})
```

```
array([[ [0., 1.], [2., 3.]],  
       [ [0., 2.], [4., 6.]]], dtype=float32)
```

⇒ Dimension will only be fully defined at runtime. (Often used for batch dimension.)

⇒ The shape information can be accessed **as a Tensor**.

# Sparse Tensors

```
sparse = tf.SparseTensor(indices=[[0, 0], [1, 2]], values=[1, 2], dense_shape=[3, 4])  
sparse.eval()
```

```
SparseTensorValue(indices=array([[0, 0],  
                                  [1, 2]]), values=array([1, 2], dtype=int32),  
dense_shape=array([3, 4]))
```

```
tf.sparse_tensor_to_dense(sparse).eval()
```

```
array([[1, 0, 0, 0],  
       [0, 0, 2, 0],  
       [0, 0, 0, 0]], dtype=int32)
```

[0,0]

[1,2]

⇒ Efficient encoding of "mostly zero" tensors (one-hot, linear models)

⇒ Variable-length features

⇒ Most Ops only for dense tensors


# Variables

```
v = tf.Variable(1, name="v")
print(v.eval())
```

FailedPreconditionError ...


```
init_op = tf.global_variables_initializer()
sess.run(init_op)
print(v.eval())
```

1




## "v++"

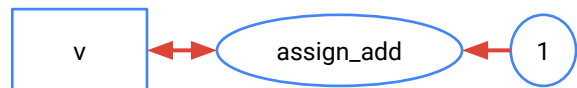
```
v = tf.Variable(0, name="v")
tf.global_variables_initializer().run()
v = v + 1
```



```
v = tf.Variable(0, name="v")
tf.global_variables_initializer().run()
sess.run([tf.assign(v, v.eval() + 1)])
sess.run([tf.assign(v, v.eval() + 1)])
```



```
v = tf.Variable(0, name="v")
tf.global_variables_initializer().run()
inc_v = tf.assign_add(v, 1)
sess.run([inc_v])
sess.run([inc_v])
```




# Naming, Sharing

```
v1 = tf.Variable(0, name="v1")
v2 = tf.get_variable("v2", initializer=0)
v3 = tf.get_variable("v3", shape=(3, 3), initializer=tf.random_normal_initializer())
v1 = tf.Variable(0, name="v1")
v2 = tf.get_variable("v2", initializer=0)
```

**ValueError:** Variable v2 already exists

```
with tf.variable_scope("vars"):
    x = tf.get_variable("x", initializer=0.)

with tf.variable_scope("vars", reuse=True):
    x = tf.get_variable("x", initializer=0.)
```




# Sharing Example

```
def component(x, n, initializer=tf.random_normal_initializer):
    with tf.variable_scope("component", reuse=tf.AUTO_REUSE):
        weights = tf.get_variable("weights", shape=(n, n), initializer=initializer)
        biases = tf.get_variable("biases", shape=(n,), initializer=initializer)
        return tf.matmul(weights, x) + biases
```

```
with tf.name_scope("tower0"):
    x = tf.random_normal([5, 1])
    y0 = component(x, n=5)
```

```
with tf.name_scope("tower1"):
    x = tf.random_normal([5, 1])
    y1 = component(x, n=5)
```

```
y = y0 + y1
```




# Collections

```
x = tf.get_variable("x", (), tf.float32)
y = tf.get_variable("y", (), tf.float32, collections=())
z = tf.get_variable("z", (), tf.float32, collections=(), trainable=False)
graph.get_collection(tf.GraphKeys.TRAINABLE_VARIABLES)
```

```
[<tf.Variable 'x:0' shape=() dtype=float32_ref>,
 <tf.Variable 'y:0' shape=() dtype=float32_ref>]
```

```
graph.get_collection(tf.GraphKeys.GLOBAL_VARIABLES)
```

```
[<tf.Variable 'x:0' shape=() dtype=float32_ref>]
```



# Save / Restore State

```
with tf.Graph().as_default() as graph, tf.Session() as sess:
    v = tf.get_variable("v", initializer=0.)
    ...
    saver = tf.train.Saver()
    save_path = saver.save(sess, "/tmp/model.ckpt")

with tf.Session(graph=graph) as sess:
    saver.restore(sess, "/tmp/model.ckpt")
    print(v.eval())
```

## Checkpoint

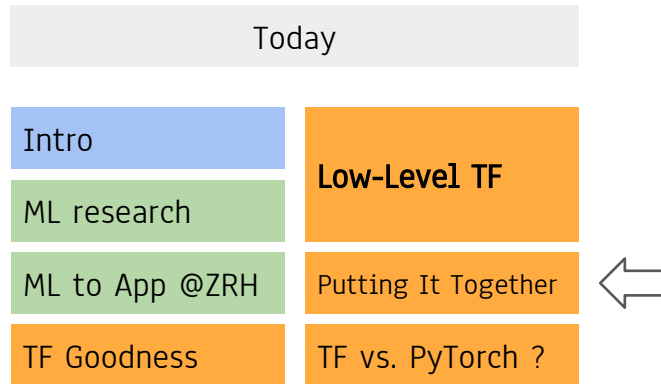
```
/tmp/model.ckpt.data-*
/tmp/model.ckpt.index
/tmp/model.ckpt.meta
```

```
from tensorflow.python.tools import inspect_checkpoint as chkp
chkp.print_tensors_in_checkpoint_file(
    "/tmp/model.ckpt", tensor_name='', all_tensors=True, all_tensor_names=True)
```

```
tensor_name: v
1.0
```



# Agenda



## Imports, ...

```
import tensorflow as tf
print(tf.__version__)
from tensorflow.python.client import device_lib
device_lib.list_local_devices()
```

```
1.7.0
[name: "/device:CPU:0"
 device_type: "CPU"
 ...
 name: "/device:GPU:0"
 device_type: "GPU"
 ...
 physical_device_desc: "device: 0, name: Tesla K80, ..."]
...
```

## ..., Data, ...

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)

mnist.train.next_batch(2)
```

```
...
Successfully downloaded train-images-idx3-ubyte.gz 9912422 bytes.
...
(array([[0., 0., 0., ..., 0., 0., 0.],
        [0., 0., 0., ..., 0., 0., 0.]], dtype=float32),
 array([[0., 0., 0., 1., 0., 0., 0., 0., 0., 0.],
        [0., 0., 0., 1., 0., 0., 0., 0., 0., 0.])))
```

## ..., Defining the Model, ...

```
graph = tf.Graph()

with graph.as_default():
    x = tf.placeholder(tf.float32, [None, 784])
    y_ = tf.placeholder(tf.float32, [None, 10])

    W = tf.Variable(tf.zeros([784, 10]))
    b = tf.Variable(tf.zeros([10]))

    logits = tf.matmul(x, W) + b
    y = tf.nn.softmax(logits)

    cross_entropy = tf.reduce_mean(-tf.reduce_sum(y_ * tf.log(y), axis=1))
    optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.5)
    train_step = optimizer.minimize(cross_entropy)
```

# ..., Train + Evaluate

..., defining the model, ...

```
x = tf.placeholder(...)
y_ = tf.placeholder(...)
...
y = tf.nn.softmax(logits)
```

```
with tf.Session(graph=graph) as sess:
    sess.run(tf.global_variables_initializer())

    # train
    for _ in range(1100):
        batch_xs, batch_ys = mnist.train.next_batch(100)
        sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})

    # evaluate
    correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    print(sess.run(accuracy, feed_dict={x: mnist.test.images, y_: mnist.test.labels}))
```

0.9183

## Agenda

Today

Intro

ML research

ML to App @ZRH

TF Goodness

**Low-Level TF**

Putting It Together

TF vs. PyTorch ?



# TensorFlow vs. PyTorch ?

	TensorFlow	PyTorch
Adoption	Huge	Rapidly growing
Graph	static, separate	dynamic, pythonic
Visualization	TensorBoard	Matplotlib etc
Deployment	TensorFlow Serving TensorFlow Lite, tensorflow.js	Caffe2
Parallelization	Distributed training device placement, variable scope	torch.nn.DataParallel
Abstractions	tf.keras, tf.layers, tf.estimator, tf.slim, ...	torch.nn.Module
	<b>BIG</b>	<b>NEW</b>

## tf.keras

```
from keras.models import Sequential
from keras.layers import Dense

model = Sequential([
    Dense(10, activation='softmax', input_dim=784)])

model.compile(loss='categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])

model.fit(mnist.train.images, mnist.train.labels, epochs=2)
model.evaluate(mnist.test.images, mnist.test.labels)
```

# TensorFlow Eager

- No Session – No Graph.
- All tensors are executed "eager"ly.
- `tf.GradientTape` records gradients for `tf.eager.Variable`

## Develop?

- ✓ Experimentation.
- ✓ Debugging, introspection.
- ✓ Natural control-flow.
- ✓ Very similar to PyTorch.

## Deploy?

- ✗ Not mature yet.
- ✗ No graph-based optimization.
- ✗ No graph-based distribution / replication.

# TensorFlow Eager – Putting it Together

```
W = tf.contrib.eager.Variable(tf.zeros([784, 10]))
b = tf.contrib.eager.Variable(tf.zeros([10]))
```

```
def get_preds(x):
    return tf.nn.softmax(tf.matmul(x, W) + b)
```

Define model

```
for _ in range(1100):
    batch_xs, batch_ys = mnist.train.next_batch(100)
    with tf.contrib.eager.GradientTape() as tape:
        preds = get_preds(batch_xs)
        loss = tf.reduce_mean(-tf.reduce_sum(batch_ys * tf.log(preds), axis=1))
        dW, db = tape.gradient(loss, [W, b])
        W.assign_sub(dW * 0.5)
        b.assign_sub(db * 0.5)
```

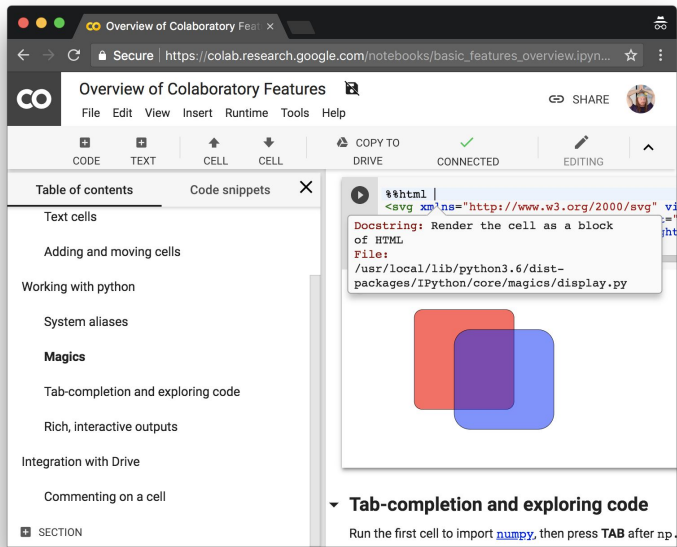
Train

```
accuracy = lambda x,y: (x.argmax(axis=1) == y.argmax(axis=1)).mean()
accuracy(get_preds(mnist.test.images).numpy(), mnist.test.labels)
```

Evaluate

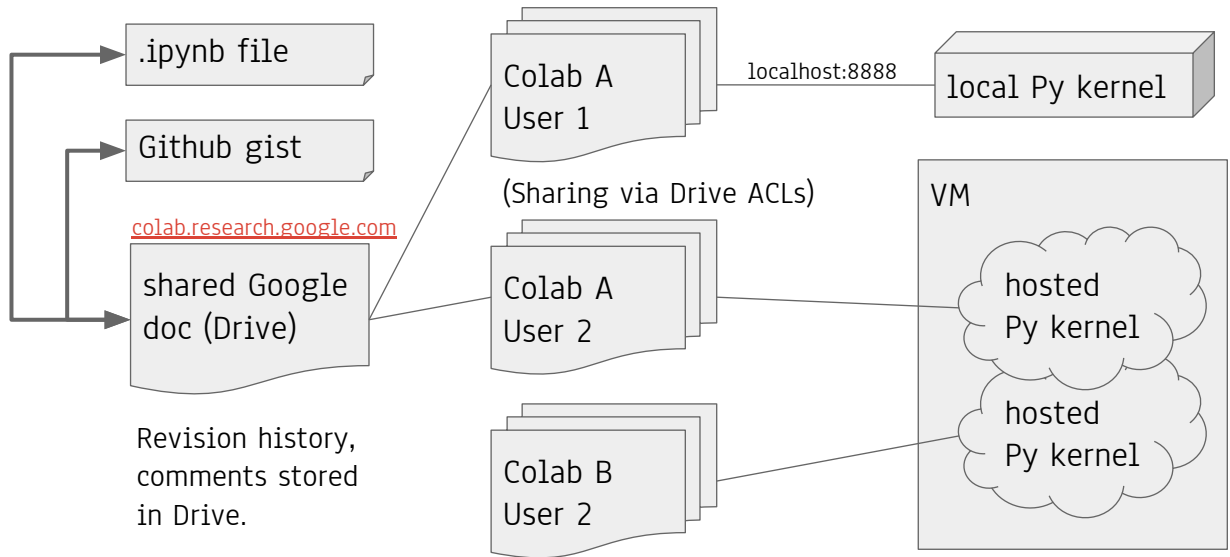
# Practical Session, Colab

## IPython ... Jupyter ... Colab



- TOC, folding
- Popover help
- Forms
- Revision history
- Document comments
- Snippets

# Colab : About Kernels



## Finding Documentation

1. Colab auto-completion

2. <https://tensorflow.org>

- a. Programmer's guide
- b. Tutorials
- c. API docs

3. Internet search : Stack Overflow, ...

4. Github codesearch : Search for usage patterns.

```
builder.add_meta_graph_and_variables([  
    sess,  
    [tf.saved_model.tag_def(sess.graph_def, sess.graph_def.tags)],  
    signature_def_map={  
        tf.saved_model.signature_def_map_key: SignatureDef(  
            inputs={'x':  
                Docstring:  
                outputs={'y':  
                method_name=  
                Signature: builder.add_meta_graph_and_variables(sess,  
                    assets_collection=None, legacy_init_op=None, clear  
                    strip_default_attrs=False)  
                Docstring:  
                Adds the current meta graph to the SavedModel and s
```

# Thanks!



<https://goo.gl/aSLGAA>

## References 1/3

(Agüera y Arcas 2017) Now Playing: Continuous low-power music recognition  
<https://arxiv.org/abs/1711.10958>

(Bello 2017) Neural Optimizer Search with Reinforcement Learning  
<https://arxiv.org/abs/1709.07417>

(Cer 2018) Universal Sentence Encoder  
<https://arxiv.org/abs/1803.11175>

(Chiu 2016) Speech recognition for medical conversations  
<https://arxiv.org/abs/1711.07274>

(Chan 2015) Listen, Attend and Spell  
<https://arxiv.org/abs/1508.01211>

(Gulshan 2016) Development and Validation of a Deep Learning Algorithm for Detection of Diabetic Retinopathy in Retinal Fundus Photographs  
<https://jamanetwork.com/journals/jama/fullarticle/2588763>



## References 2/3

(Jouppi 2017) In-Datacenter Performance Analysis of a Tensor Processing Unit

<https://arxiv.org/abs/1704.04760>

(Liu 2017) Detecting Cancer Metastases on Gigapixel Pathology Images

<https://arxiv.org/abs/1703.02442>

(Po-Hsuan 2018) An Augmented Reality Microscope for Real-time Automated Detection of Cancer  
*Under Review*

(Poplin 2016) Creating a universal SNP and small indel variant caller with deep neural networks

<https://www.biorxiv.org/content/early/2016/12/21/092890>

(Rajkumar 2018) Scalable and accurate deep learning for electronic health records

<https://arxiv.org/abs/1801.07860>

(Ramachandran 2017) Searching for Activation Functions

<https://arxiv.org/abs/1710.05941>

## References 3/3

(Shen 2017) Natural TTS Synthesis by Conditioning WaveNet on Mel Spectrogram Predictions

<https://arxiv.org/abs/1712.05884>

(Van den Oord 2016) WaveNet: A Generative Model for Raw Audio

<https://arxiv.org/abs/1609.03499>

(Zoph 2016) Neural Architecture Search with Reinforcement Learning

<https://arxiv.org/abs/1611.01578>

(Zoph 2017) Learning Transferable Architectures for Scalable Image Recognition

<https://arxiv.org/abs/1707.07012>

# Links 1/2

Introduction to TensorFlow Lite

<https://www.tensorflow.org/mobile/tflite/>

Machine Learning in JavaScript (TensorFlow Dev Summit 2018)

<https://youtu.be/YB-kfeNIPCE>

ImageNet is the new MNIST

<https://supercomputersfordl2017.github.io/Presentations/ImageNetNewMNIST.pdf>

Distributed TensorFlow

<https://www.tensorflow.org/deploy/distributed>

Distributed TensorFlow (TensorFlow Dev Summit 2017)

[https://youtu.be/la\\_M6bCV91M](https://youtu.be/la_M6bCV91M)

Graphs and Sessions

[https://www.tensorflow.org/programmers\\_guide/graphs](https://www.tensorflow.org/programmers_guide/graphs)

# Links 2/2

Sparse Tensors

[https://www.tensorflow.org/versions/master/api\\_guides/python/sparse\\_ops](https://www.tensorflow.org/versions/master/api_guides/python/sparse_ops)

TensorFlow Linear Model Tutorial

<https://www.tensorflow.org/versions/master/tutorials/wide>

Variables

[https://www.tensorflow.org/programmers\\_guide/variables](https://www.tensorflow.org/programmers_guide/variables)

Variable Ops

[https://www.tensorflow.org/versions/master/api\\_guides/python/state\\_ops#Variables](https://www.tensorflow.org/versions/master/api_guides/python/state_ops#Variables)

Sharing Variables

[https://www.tensorflow.org/versions/master/programmers\\_guide/variable\\_scope](https://www.tensorflow.org/versions/master/programmers_guide/variable_scope)

Eager Execution

[https://www.tensorflow.org/programmers\\_guide/eager](https://www.tensorflow.org/programmers_guide/eager)